

Carrie Final Project Data Analysis_2000

November 24, 2024

1 SIPRI Military Expenditure dataset

1.1 Carrie Little

1.1.1 Import Necessary Libraries

```
[1]: # Import All Necessary Libraries
import numpy as np
import pandas as pd
from pathlib import Path
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import cvxpy as cp
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
[2]: # Load Dataset
from pathlib import Path
import pandas as pd

# Define the file path
file_path = Path("..") / "data" /
    ↪ "Cleaned_merged_SIPRI_Region_ACLED_starting2000.csv"      # Suitable for
    ↪ within Github repository
# file_path = 'Cleaned_merged_SIPRI_Region_ACLED_starting2000.csv'
    ↪ # for when working in same directory

# Load the CSV file
df = pd.read_csv(file_path)

df.head()
```

```
[2]: Country Year Expenditure-Share_of_Govt_spending Expenditure-Share_of_GDP \
0 Brazil 2000 NaN 0.017307
1 Brazil 2001 0.047167 0.019519
```

2	Brazil	2002	0.041112	0.018958
3	Brazil	2003	0.035175	0.015035
4	Brazil	2004	0.035561	0.014613

	Expenditure_Per_Capita	Expenditure_Constant_2022	Region	Subregion \
0	64.500907	15643.43168	Americas	South America
1	61.332747	18119.43481	Americas	South America
2	53.550103	18361.33076	Americas	South America
3	45.956041	14647.82954	Americas	South America
4	52.945029	15219.74744	Americas	South America

	index_level	total_rank_from_avg_rank	...	total_score_rank \
0	1	6	...	6
1	1	6	...	6
2	1	6	...	6
3	1	6	...	6
4	1	6	...	6

	Deadliness_raw	Diffusion_raw	Danger_raw	Fragmentation_raw \
0	6678	0.003	4117	67
1	6678	0.003	4117	67
2	6678	0.003	4117	67
3	6678	0.003	4117	67
4	6678	0.003	4117	67

	Deadliness_scaled	Diffusion_scaled	Danger_scaled	Fragmentation_scaled \
0	0.167247	0.004342	0.635536	0.044108
1	0.167247	0.004342	0.635536	0.044108
2	0.167247	0.004342	0.635536	0.044108
3	0.167247	0.004342	0.635536	0.044108
4	0.167247	0.004342	0.635536	0.044108

	total_score
0	0.851
1	0.851
2	0.851
3	0.851
4	0.851

[5 rows x 29 columns]

1.1.2 Descriptive statistics

```
[3]: # Descriptive Information
      # Columns
      columns = df.info()
      ↪ Create a Dataframe of Descriptive Statistics
```

#

columns

#

↪ *Display the Dataframe of Descriptive Statistics*

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3716 entries, 0 to 3715
```

```
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	Country	3716 non-null	object
1	Year	3716 non-null	int64
2	Expenditure-Share_of_Govt_spending	3533 non-null	float64
3	Expenditure-Share_of_GDP	3602 non-null	float64
4	Expenditure_Per_Capita	3604 non-null	float64
5	Expenditure_Constant_2022	3673 non-null	float64
6	Region	3716 non-null	object
7	Subregion	3716 non-null	object
8	index_level	3716 non-null	int64
9	total_rank_from_avg_rank	3716 non-null	int64
10	avg_rank	3716 non-null	float64
11	Deadliness_rank	3716 non-null	int64
12	Diffusion_rank	3716 non-null	int64
13	Danger_rank	3716 non-null	int64
14	Fragmentation_rank	3716 non-null	int64
15	Deadliness_scaled_rank	3716 non-null	int64
16	Diffusion_scaled_rank	3716 non-null	int64
17	Danger_scaled_rank	3716 non-null	int64
18	Fragmentation_scaled_rank	3716 non-null	int64
19	total_score_rank	3716 non-null	int64
20	Deadliness_raw	3716 non-null	int64
21	Diffusion_raw	3716 non-null	float64
22	Danger_raw	3716 non-null	int64
23	Fragmentation_raw	3716 non-null	int64
24	Deadliness_scaled	3716 non-null	float64
25	Diffusion_scaled	3716 non-null	float64
26	Danger_scaled	3716 non-null	float64
27	Fragmentation_scaled	3716 non-null	float64
28	total_score	3716 non-null	float64

```
dtypes: float64(11), int64(15), object(3)
```

```
memory usage: 842.0+ KB
```

[4]: # *Descriptive Statistics*

Describe

```
summary_stats = df.describe()
```

↪ # *Create a Dataframe of Descriptive Statistics*

```
summary_stats
```

↪ # *Display the Dataframe of Descriptive Statistics*

```
[4]:      Year Expenditure-Share_of_Govt_spending \
count  3716.000000      3533.000000
mean   2011.556512      0.071187
std     6.889732      0.060611
min     2000.000000      0.000672
25%     2006.000000      0.031740
50%     2012.000000      0.051833
75%     2018.000000      0.091994
max     2023.000000      0.581707
```

```
      Expenditure-Share_of_GDP Expenditure_Per_Capita \
count      3602.000000      3604.000000
mean         0.019767      243.374084
std         0.018137      426.197550
min         0.000163       0.071803
25%         0.010229      17.077744
50%         0.014996      69.809759
75%         0.023754      270.439168
max         0.366531      5718.771025
```

```
      Expenditure_Constant_2022 index_level total_rank_from_avg_rank \
count      3673.000000  3716.000000      3716.000000
mean      11760.889052    3.491119      86.736006
std      68296.745904    0.894834      49.421439
min         0.000000    1.000000       2.000000
25%       144.314628    3.000000      45.000000
50%       739.329049    4.000000      85.000000
75%      4523.330790    4.000000     131.000000
max     990485.412100    4.000000     160.000000
```

```
      avg_rank Deadliness_rank Diffusion_rank ... total_score_rank \
count  3716.000000      3716.000000      3716.000000 ...      3716.000000
mean    71.252758      79.672497      43.389397 ...      80.705867
std     33.403444      40.775882      13.511428 ...      42.837423
min       4.250000       2.000000       2.000000 ...       2.000000
25%     42.750000      43.000000      37.000000 ...      45.000000
50%     76.500000      84.000000      51.000000 ...      81.000000
75%    101.000000     125.000000      51.000000 ...     117.000000
max    114.750000     125.000000      51.000000 ...     136.000000
```

```
      Deadliness_raw Diffusion_raw Danger_raw Fragmentation_raw \
count  3716.000000      3716.000000  3716.000000      3716.000000
mean    918.312971       0.005629    252.864101      17.972282
std    3629.622878       0.023324    710.539730     107.454813
min         0.000000       0.000000       0.000000       0.000000
25%         0.000000       0.000000       2.000000       0.000000
50%         7.000000       0.000000      13.000000       2.000000
```

75%	246.000000	0.001000	152.000000	9.000000
max	37014.000000	0.176000	5426.000000	1519.000000

	Deadliness_scaled	Diffusion_scaled	Danger_scaled \
count	3716.000000	3716.000000	3716.000000
mean	0.022999	0.008146	0.039034
std	0.090902	0.033754	0.109685
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000309
50%	0.000175	0.000000	0.002007
75%	0.006161	0.001447	0.023464
max	0.926995	0.254703	0.837604

	Fragmentation_scaled	total_score
count	3716.000000	3716.000000
mean	0.011832	0.081966
std	0.070740	0.231303
min	0.000000	0.000000
25%	0.000000	0.001000
50%	0.001317	0.005000
75%	0.005925	0.037000
max	1.000000	1.940000

[8 rows x 26 columns]

1.1.3 Histogram

```
[5]: # Data Preprocessing and Summary Statistics

# Checking for missing values
missing_values = df.isnull().sum()

# Summary statistics for numerical data
summary_statistics = df.describe()

# Viewing the columns for a better understanding of categorical vs numerical
↳ data
df_columns = df.columns

# Plotting histograms for numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
df[numerical_columns].hist(figsize=(15, 15), bins=20)
plt.suptitle('Histograms of Numerical Data', size=16)
plt.show()

# Plotting bar graphs for categorical columns
categorical_columns = df.select_dtypes(include=['object']).columns
```

```
plt.figure(figsize=(15, 20))
for i, column in enumerate(categorical_columns):
    plt.subplot(len(categorical_columns), 1, i + 1)
    sns.countplot(data=df, x=column)
    plt.xticks(rotation=45)
    plt.title(f'Bar Graph of {column}')

plt.tight_layout()
plt.show()

# Display summary statistics
summary_statistics

# Displaying missing values information
missing_values
```

Histograms of Numerical Data

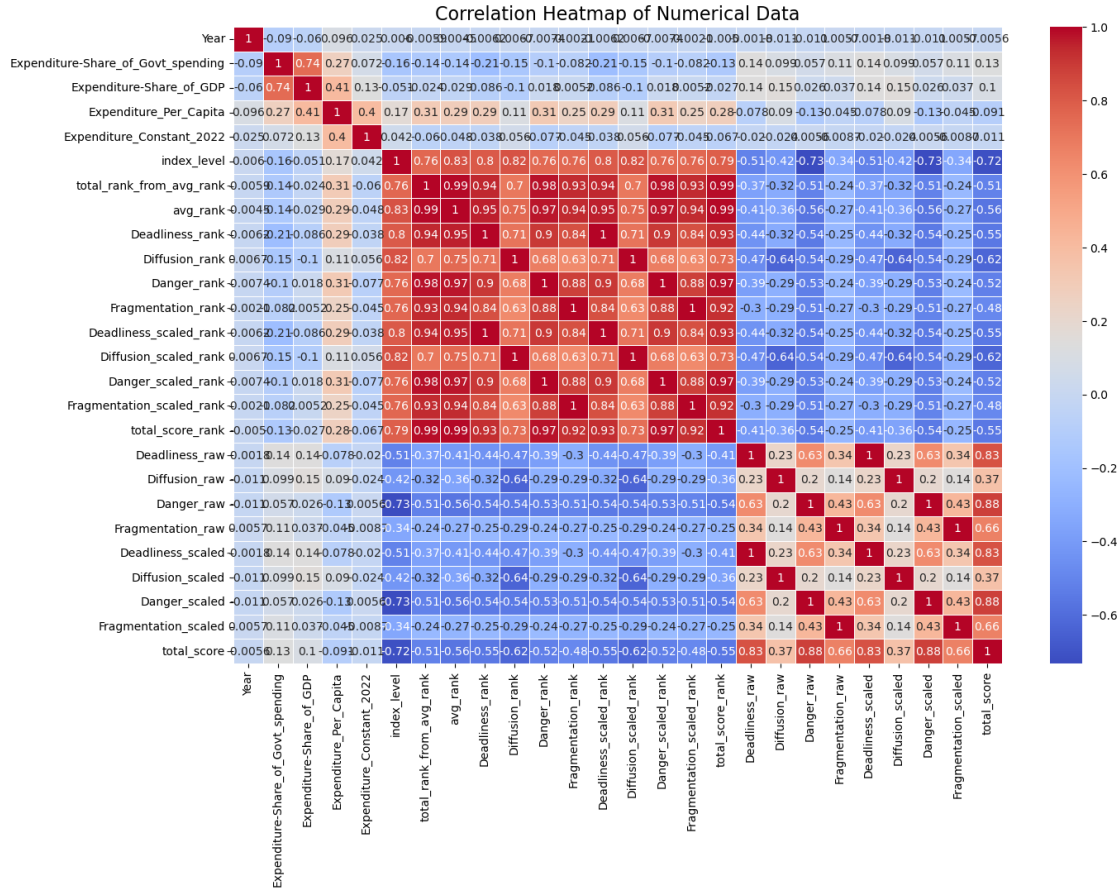


Expenditure-Share_of_Govt_spending	183
Expenditure-Share_of_GDP	114
Expenditure_Per_Capita	112
Expenditure_Constant_2022	43
Region	0
Subregion	0
index_level	0
total_rank_from_avg_rank	0
avg_rank	0
Deadliness_rank	0
Diffusion_rank	0
Danger_rank	0
Fragmentation_rank	0
Deadliness_scaled_rank	0
Diffusion_scaled_rank	0
Danger_scaled_rank	0
Fragmentation_scaled_rank	0
total_score_rank	0
Deadliness_raw	0
Diffusion_raw	0
Danger_raw	0
Fragmentation_raw	0
Deadliness_scaled	0
Diffusion_scaled	0
Danger_scaled	0
Fragmentation_scaled	0
total_score	0
dtype: int64	

```
[6]: # Generating a correlation heatmap for numerical data

# Calculate the correlation matrix for numerical columns
correlation_matrix = df[numerical_columns].corr()

# Plotting the heatmap using seaborn
plt.figure(figsize=(15, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap of Numerical Data', size=16)
plt.show()
```



```
[7]: # Identifying highly correlated features (absolute correlation > 0.7)
highly_correlated_pairs = correlation_matrix.abs().unstack().
    ↪sort_values(ascending=False)
highly_correlated_pairs = highly_correlated_pairs[highly_correlated_pairs < 1.
    ↪0] # Exclude self-correlation

# Get the top correlated pairs
top_correlated_pairs = highly_correlated_pairs[highly_correlated_pairs > 0.7].
    ↪index

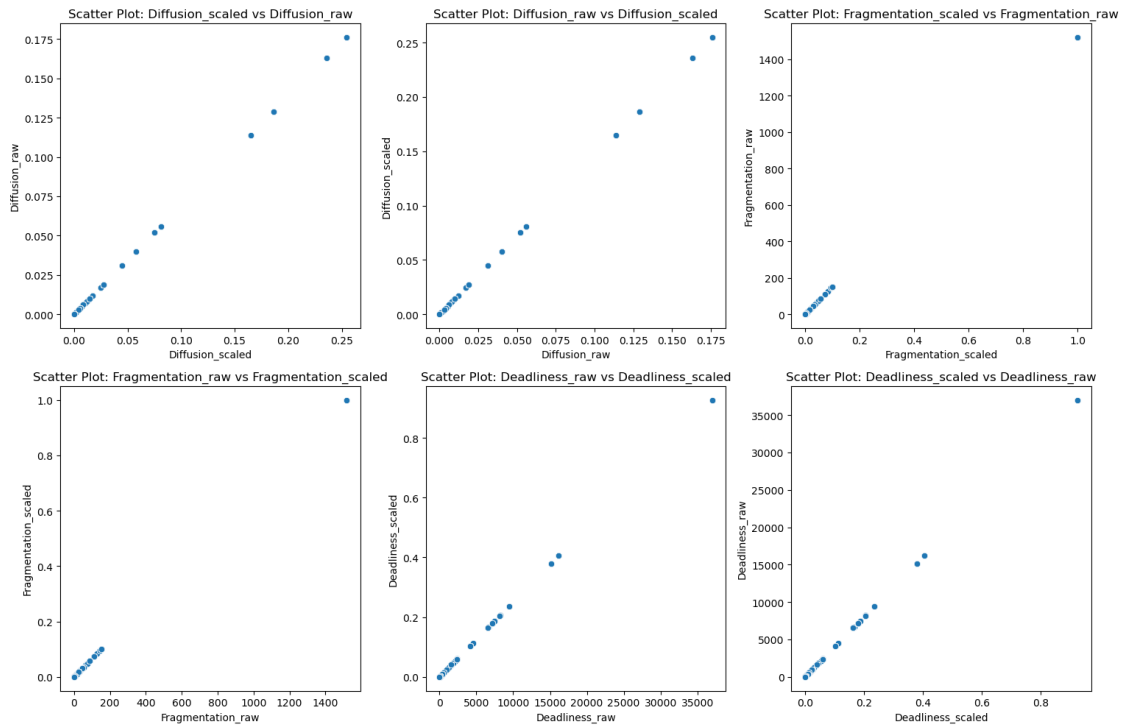
# Plotting scatter plots for highly correlated feature pairs
plt.figure(figsize=(15, 10))
plot_index = 1
for (feature1, feature2) in top_correlated_pairs:
    if plot_index > 6: # Limit the number of scatter plots to display
        break
    plt.subplot(2, 3, plot_index)
    sns.scatterplot(data=df, x=feature1, y=feature2)
    plt.title(f'Scatter Plot: {feature1} vs {feature2}')
```

```

plot_index += 1

plt.tight_layout()
plt.show()

```



```

[8]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

```

```

[9]: # Load Dataset
from pathlib import Path
import pandas as pd

# Define the file path
file_path = Path("..") / "data" /
    ↳ "Cleaned_merged_SIPRI_Region_ACLED_starting2000.csv"      # Suitable for
    ↳ within Github repository
# file_path = 'Cleaned_merged_SIPRI_Region_ACLED_starting2000.csv'
    ↳ # for when working in same directory

```

```
# Load the CSV file
data = pd.read_csv(file_path)

data.head()
```

```
[9]: Country Year Expenditure-Share_of_Govt_spending Expenditure-Share_of_GDP \
0 Brazil 2000 NaN 0.017307
1 Brazil 2001 0.047167 0.019519
2 Brazil 2002 0.041112 0.018958
3 Brazil 2003 0.035175 0.015035
4 Brazil 2004 0.035561 0.014613
```

```
Expenditure_Per_Capita Expenditure_Constant_2022 Region Subregion \
0 64.500907 15643.43168 Americas South America
1 61.332747 18119.43481 Americas South America
2 53.550103 18361.33076 Americas South America
3 45.956041 14647.82954 Americas South America
4 52.945029 15219.74744 Americas South America
```

```
index_level total_rank_from_avg_rank ... total_score_rank \
0 1 6 ... 6
1 1 6 ... 6
2 1 6 ... 6
3 1 6 ... 6
4 1 6 ... 6
```

```
Deadliness_raw Diffusion_raw Danger_raw Fragmentation_raw \
0 6678 0.003 4117 67
1 6678 0.003 4117 67
2 6678 0.003 4117 67
3 6678 0.003 4117 67
4 6678 0.003 4117 67
```

```
Deadliness_scaled Diffusion_scaled Danger_scaled Fragmentation_scaled \
0 0.167247 0.004342 0.635536 0.044108
1 0.167247 0.004342 0.635536 0.044108
2 0.167247 0.004342 0.635536 0.044108
3 0.167247 0.004342 0.635536 0.044108
4 0.167247 0.004342 0.635536 0.044108
```

```
total_score
0 0.851
1 0.851
2 0.851
3 0.851
4 0.851
```

[5 rows x 29 columns]

1.1.4 Data Cleaning for Regression Models

```
[10]: # Drop the columns that are not relevant for the forecasting model as mentioned
      ↪by the user
columns_to_drop = [
    'Deadliness_scaled', 'Diffusion_scaled', 'Danger_scaled',
    ↪'Fragmentation_scaled',
    'Deadliness_rank', 'Diffusion_rank', 'Danger_rank', 'Fragmentation_rank'
]
data_cleaned = data.drop(columns=columns_to_drop, errors='ignore')

# Display the cleaned dataset to check the changes
data_cleaned.head()
```

```
[10]: Country Year Expenditure-Share_of_Govt_spending Expenditure-Share_of_GDP \
0 Brazil 2000 NaN 0.017307
1 Brazil 2001 0.047167 0.019519
2 Brazil 2002 0.041112 0.018958
3 Brazil 2003 0.035175 0.015035
4 Brazil 2004 0.035561 0.014613
```

```
Expenditure_Per_Capita Expenditure_Constant_2022 Region Subregion \
0 64.500907 15643.43168 Americas South America
1 61.332747 18119.43481 Americas South America
2 53.550103 18361.33076 Americas South America
3 45.956041 14647.82954 Americas South America
4 52.945029 15219.74744 Americas South America
```

```
index_level total_rank_from_avg_rank ... Deadliness_scaled_rank \
0 1 6 ... 10
1 1 6 ... 10
2 1 6 ... 10
3 1 6 ... 10
4 1 6 ... 10
```

```
Diffusion_scaled_rank Danger_scaled_rank Fragmentation_scaled_rank \
0 31 3 9
1 31 3 9
2 31 3 9
3 31 3 9
4 31 3 9
```

```
total_score_rank Deadliness_raw Diffusion_raw Danger_raw \
0 6 6678 0.003 4117
1 6 6678 0.003 4117
```

2	6	6678	0.003	4117
3	6	6678	0.003	4117
4	6	6678	0.003	4117

	Fragmentation_raw	total_score
0	67	0.851
1	67	0.851
2	67	0.851
3	67	0.851
4	67	0.851

[5 rows x 21 columns]

```
[11]: # Check for missing values in the cleaned dataset
missing_values = data_cleaned.isna().sum()

# Display the columns with missing values and the count of missing values
missing_values[missing_values > 0]
```

```
[11]: Expenditure-Share_of_Govt_spending    183
Expenditure-Share_of_GDP                  114
Expenditure_Per_Capita                    112
Expenditure_Constant_2022                  43
dtype: int64
```

```
[12]: # Create a copy of the dataframe to avoid modifying the original data
data_cleaned = data_cleaned.copy()

# Handling missing values in "Expenditure_Constant_2022" using forward and
↳backward filling
data_cleaned['Expenditure_Constant_2022'] =
↳data_cleaned['Expenditure_Constant_2022'].ffill().bfill()

# Handling missing values in other numerical columns using median imputation
for column in ['Expenditure-Share_of_Govt_spending',
↳'Expenditure-Share_of_GDP', 'Expenditure_Per_Capita']:
    median_value = data_cleaned[column].median()
    data_cleaned[column] = data_cleaned[column].fillna(median_value)

# Verify that there are no missing values left in the dataset
missing_values_after_imputation = data_cleaned.isna().sum()
missing_values_after_imputation[missing_values_after_imputation > 0]
```

```
[12]: Series([], dtype: int64)
```

Supervised Learning (Regression)

1.1.5 Linear Regression Model 1

```
[13]: # Selecting relevant features for forecasting
# We will primarily use 'Year' and 'Expenditure_Constant_2022' for the
↳time-series analysis.
forecast_data = data_cleaned[['Country', 'Year', 'Expenditure_Constant_2022']]

# Sorting the data by Country and Year to prepare for time-series modeling
forecast_data = forecast_data.sort_values(by=['Country', 'Year'])

# Display the prepared data for forecasting
forecast_data.head()
```

```
[13]:      Country  Year  Expenditure_Constant_2022
619  Afghanistan  2004             183.648324
620  Afghanistan  2005             165.378054
621  Afghanistan  2006             167.189480
622  Afghanistan  2007             257.366345
623  Afghanistan  2008             224.292155
```

```
[14]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets for each country
# We will create a dictionary to hold the train-test split for each country
countries = forecast_data['Country'].unique()
train_test_data = {}

for country in countries:
    country_data = forecast_data[forecast_data['Country'] == country]
    X = country_data[['Year']]
    y = country_data['Expenditure_Constant_2022']

    # Split the data (80% training, 20% testing)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳shuffle=False)

    train_test_data[country] = {
        'X_train': X_train,
        'X_test': X_test,
        'y_train': y_train,
        'y_test': y_test
    }

# Display the train-test split information for a sample country
sample_country = countries[0]
train_test_data[sample_country]
```

```

[14]: {'X_train':      Year
      619  2004
      620  2005
      621  2006
      622  2007
      623  2008
      624  2009
      625  2010
      626  2011
      627  2012
      628  2013
      629  2014
      630  2015
      631  2016
      632  2017,
      'X_test':      Year
      633  2018
      634  2019
      635  2020
      636  2021,
      'y_train': 619      183.648324
      620      165.378054
      621      167.189480
      622      257.366345
      623      224.292155
      624      252.407949
      625      269.912673
      626      265.485040
      627      198.955803
      628      183.421702
      629      223.721970
      630      178.913871
      631      177.240056
      632      174.272685
      Name: Expenditure_Constant_2022, dtype: float64,
      'y_test': 633      190.335314
      634      228.106548
      635      263.904202
      636      268.600635
      Name: Expenditure_Constant_2022, dtype: float64}

```

```

[15]: from sklearn.linear_model import LinearRegression
import numpy as np

# Initialize a dictionary to store the models for each country
models = {}

```



```

# Train a Linear Regression model for each country
for country in countries:
    # Extract train-test data for the current country
    X_train = train_test_data[country]['X_train']
    y_train = train_test_data[country]['y_train']

    # Initialize and train the Linear Regression model
    model = LinearRegression()
    model.fit(X_train, y_train)

    # Store the trained model
    models[country] = model

# Forecast the next 5 and 10 years for each country using the trained models
forecast_years_5 = pd.DataFrame({'Year': np.arange(2022, 2027)})
forecast_years_10 = pd.DataFrame({'Year': np.arange(2022, 2032)})

# Create a dictionary to store forecasts for each country
forecasts = {}

for country in countries:
    # Use the trained model to make predictions for the next 5 and 10 years
    model = models[country]
    forecast_5 = model.predict(forecast_years_5)
    forecast_10 = model.predict(forecast_years_10)

    forecasts[country] = {
        '5_year_forecast': forecast_5,
        '10_year_forecast': forecast_10
    }

# Display the forecast for the first country as a sample
forecasts[sample_country]

```

```

[15]: {'5_year_forecast': array([196.33102894, 195.25294383, 194.17485873,
    193.09677362,
    192.01868852]),
      '10_year_forecast': array([196.33102894, 195.25294383, 194.17485873,
    193.09677362,
    192.01868852, 190.94060342, 189.86251831, 188.78443321,
    187.70634811, 186.628263  ])}

```

```

[45]: # Get the most recent "Expenditure_Constant_2022" for each country
latest_expenditure = forecast_data.groupby('Country', group_keys=False).apply(
    lambda x: x.loc[x['Year'] == x['Year'].max(), 'Expenditure_Constant_2022'].
    ↪values[0]
)

```

```

# Sort and select the top 15 countries
top_15_countries = latest_expenditure.sort_values(ascending=False).head(15).
↳index

# Plotting the 5-year forecast for the top 15 countries
plt.figure(figsize=(15, 10))
for country in top_15_countries:
    years = np.arange(2022, 2027)
    expenditure_forecast = forecasts[country]['5_year_forecast']
    plt.plot(years, expenditure_forecast, label=country)

plt.xlabel('Year')
plt.ylabel('Expenditure_Constant_2022')
plt.title('5-Year Forecast for the Top 15 Countries by_
↳Expenditure_Constant_2022')
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
↳3, 1), fontsize='small', ncol=1)
plt.grid(True, which="both", linestyle='--')
plt.show()

# Plotting the 10-year forecast for the top 15 countries
plt.figure(figsize=(15, 10))
for country in top_15_countries:
    years = np.arange(2022, 2032)
    expenditure_forecast = forecasts[country]['10_year_forecast']
    plt.plot(years, expenditure_forecast, label=country)

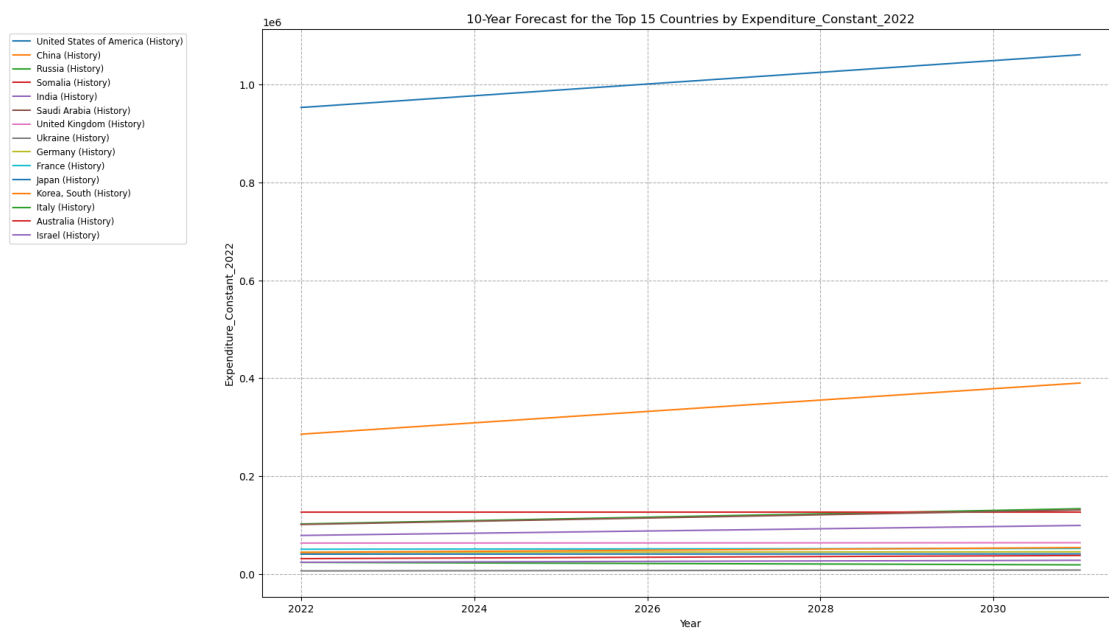
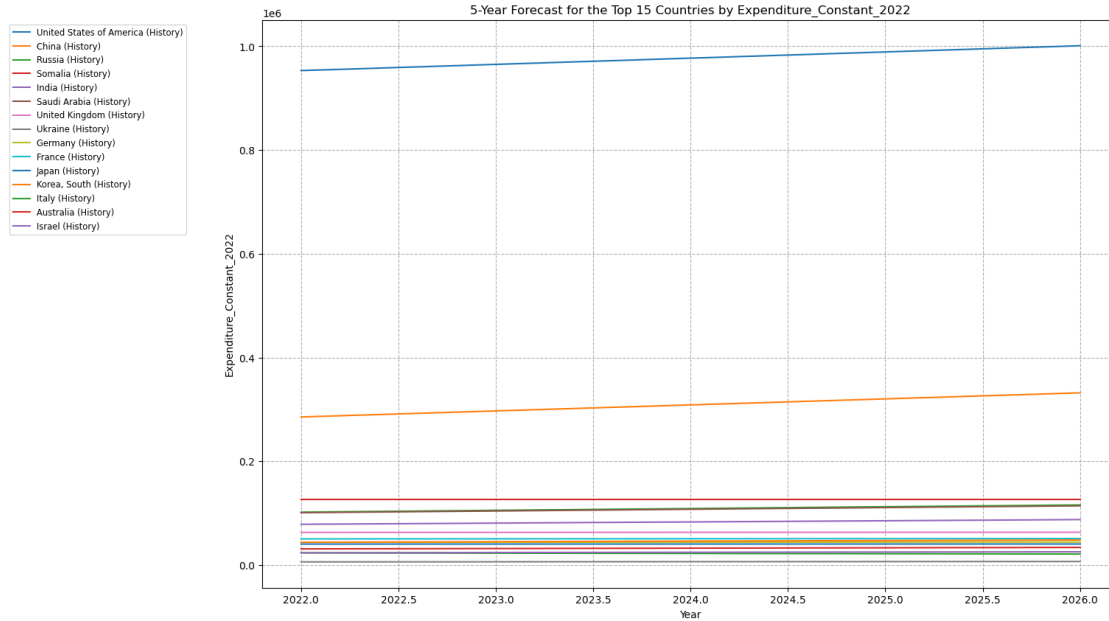
plt.xlabel('Year')
plt.ylabel('Expenditure_Constant_2022')
plt.title('10-Year Forecast for the Top 15 Countries by_
↳Expenditure_Constant_2022')
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
↳3, 1), fontsize='small', ncol=1)
plt.grid(True, which="both", linestyle='--')
plt.show()

```

C:\Users\carri\AppData\Local\Temp\ipykernel_33844\1077236069.py:2:

DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

```
latest_expenditure = forecast_data.groupby('Country', group_keys=False).apply(
```



```
[46]: # Sort the top 15 countries by the most recent expenditure level for legend
      ↪ ranking
latest_expenditure_sorted = latest_expenditure[top_15_countries].
      ↪ sort_values(ascending=False)

plt.figure(figsize=(15, 10))
```

```

country_lines = []

for country in latest_expenditure_sorted.index:
    # Get the historical data for the last 10 years for the selected country
    country_data = forecast_data[forecast_data['Country'] == country]
    historical_years = country_data['Year'].tail(10)
    historical_expenditure = country_data['Expenditure_Constant_2022'].tail(10)

    # Plot the historical data with a solid line
    historical_line, = plt.plot(
        historical_years, historical_expenditure, linestyle='-',
        label=f'{country} (History)'
    )

    # Plot the forecast for the next 10 years with a dashed line, using the
    same color
    plt.plot(
        np.arange(2022, 2032), forecasts[country]['10_year_forecast'],
        linestyle='--', color=historical_line.get_color(), label=f'{country}'
        (Forecast)
    )


    # Store reference for legend sorting
    country_lines.append((country, historical_line))

# Sort legend entries by the latest expenditure level
sorted_legend_handles = [line for country, line in sorted(country_lines,
    key=lambda x: latest_expenditure_sorted[country], reverse=True)]

plt.xlabel('Year')
plt.ylabel('Expenditure_Constant_2022')
plt.title('Historical and 10-Year Forecast for the Top 15 Countries by
    Expenditure_Constant_2022')

# Placing the legend outside to the left of the plot and ranked by most recent
expenditure level
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
    3, 1), fontsize='small', ncol=1)
plt.grid(True, which="both", linestyle='--')
plt.show()

```



```
[18]: # Displaying the formulas (coefficients and intercepts) of the Linear
      ↪ Regression models for each country

      # Store the forecast formulas for each of the top 15 countries
      forecast_formulas = {}

      for country in top_15_countries:
          # Extract the model for the current country
          model = models[country]

          # Get the coefficient (slope) and intercept of the linear regression model
          coefficient = model.coef_[0]
          intercept = model.intercept_

          # Store the formula
          forecast_formulas[country] = f"Expenditure_Constant_2022 = {coefficient:.
          ↪ 2f} * Year + {intercept:.2f}"

      # Display the forecast formulas for the top 15 countries
      forecast_formulas
```

```
[18]: {'United States of America': 'Expenditure_Constant_2022 = 11981.00 * Year +
-23272546.56',
      'China': 'Expenditure_Constant_2022 = 11618.69 * Year + -23207588.84',
      'Russia': 'Expenditure_Constant_2022 = 3464.69 * Year + -6903680.88',
      'Somalia': 'Expenditure_Constant_2022 = 0.00 * Year + 126473.35',
```

```

'India': 'Expenditure_Constant_2022 = 2265.94 * Year + -4503217.61',
'Saudi Arabia': 'Expenditure_Constant_2022 = 3301.45 * Year + -6574677.22',
'United Kingdom': 'Expenditure_Constant_2022 = 81.02 * Year + -100966.65',
'Ukraine': 'Expenditure_Constant_2022 = 194.25 * Year + -386780.27',
'Germany': 'Expenditure_Constant_2022 = 135.10 * Year + -229344.53',
'France': 'Expenditure_Constant_2022 = 174.92 * Year + -303368.71',
'Japan': 'Expenditure_Constant_2022 = 43.12 * Year + -46864.64',
'Korea, South': 'Expenditure_Constant_2022 = 1086.85 * Year + -2153701.90',
'Italy': 'Expenditure_Constant_2022 = -558.10 * Year + 1151893.60',
'Australia': 'Expenditure_Constant_2022 = 719.24 * Year + -1423274.48',
'Israel': 'Expenditure_Constant_2022 = 454.85 * Year + -896067.90'}

```

```

[19]: from sklearn.metrics import r2_score

# Prepare a list to store the country, coefficient, intercept, and R^2 values
model_summary = []

# Calculate the R^2 score for each of the top 15 countries
for country in top_15_countries:
    # Extract model and train-test data for the current country
    model = models[country]
    X_train = train_test_data[country]['X_train']
    y_train = train_test_data[country]['y_train']
    X_test = train_test_data[country]['X_test']
    y_test = train_test_data[country]['y_test']

    # Get the coefficient and intercept
    coefficient = model.coef_[0]
    intercept = model.intercept_

    # Calculate the R^2 score on the test data
    y_pred = model.predict(X_test)
    r_squared = r2_score(y_test, y_pred)

    # Append the summary to the list
    model_summary.append({
        'Country': country,
        'Coefficient': coefficient,
        'Intercept': intercept,
        'R^2 Score': r_squared
    })

# Create a DataFrame to display the model summary
model_summary_df = pd.DataFrame(model_summary)

# Display the model summary
model_summary_df

```

```
[19]:
```

	Country	Coefficient	Intercept	R ² Score
0	United States of America	11981.001032	-2.327255e+07	-25.457937
1	China	11618.685787	-2.320759e+07	0.694556
2	Russia	3464.692816	-6.903681e+06	0.290156
3	Somalia	0.000000	1.264734e+05	1.000000
4	India	2265.937042	-4.503218e+06	-0.237510
5	Saudi Arabia	3301.454063	-6.574677e+06	-78.910826
6	United Kingdom	81.022005	-1.009666e+05	-0.080564
7	Ukraine	194.253952	-3.867803e+05	-0.674017
8	Germany	135.101688	-2.293445e+05	-11.194959
9	France	174.924374	-3.033687e+05	-1.866650
10	Japan	43.120633	-4.686464e+04	-2.388573
11	Korea, South	1086.851126	-2.153702e+06	-3.369126
12	Italy	-558.095940	1.151894e+06	-10.028473
13	Australia	719.242676	-1.423274e+06	0.406479
14	Israel	454.845309	-8.960679e+05	-0.070533

1.1.6 Random Forest Regressor

```
[20]: from sklearn.ensemble import RandomForestRegressor

# Initialize a dictionary to store the Random Forest models for each country
rf_models = {}

# Train a Random Forest Regressor for each country
for country in top_15_countries:
    # Extract train-test data for the current country
    X_train = train_test_data[country]['X_train']
    y_train = train_test_data[country]['y_train']

    # Initialize and train the Random Forest Regressor
    rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
    rf_model.fit(X_train, y_train)

    # Store the trained model
    rf_models[country] = rf_model

# Forecast the next 10 years for each country using the trained Random Forest
# models
rf_forecasts = {}

for country in top_15_countries:
    # Use the trained model to make predictions for the next 10 years
    rf_model = rf_models[country]
    forecast_10 = rf_model.predict(forecast_years_10)

    rf_forecasts[country] = {
```

```

        '10_year_forecast': forecast_10
    }

```

```

# Display the forecast for one of the top 15 countries

```

```

sample_country = top_15_countries[0]
rf_forecasts[sample_country]

```

```

[20]: {'10_year_forecast': array([789481.619213, 789481.619213, 789481.619213,
789481.619213,
789481.619213, 789481.619213, 789481.619213, 789481.619213,
789481.619213, 789481.619213])}

```

```

[47]: # Plotting the historical data (last 10 years) and the 10-year Random Forest
      ↪forecast for the top 15 countries with consistent colors

```

```

plt.figure(figsize=(15, 10))

```

```

for country in sorted(top_15_countries):
    # Get the historical data for the last 10 years for the selected country
    country_data = forecast_data[forecast_data['Country'] == country]
    historical_years = country_data['Year'].tail(10)
    historical_expenditure = country_data['Expenditure_Constant_2022'].tail(10)

```

```

    # Plot the historical data with a solid line

```

```

    historical_line, = plt.plot(
        historical_years, historical_expenditure, linestyle='-',
        ↪label=f'{country} (History)'
    )

```

```

    # Plot the Random Forest forecast for the next 10 years with a dashed line,
    ↪using the same color

```

```

    forecast_years = np.arange(2022, 2032)
    expenditure_forecast = rf_forecasts[country]['10_year_forecast']
    plt.plot(
        forecast_years, expenditure_forecast, linestyle='--',
        ↪color=historical_line.get_color(), label=f'{country} (RF Forecast)'
    )

```

```

# Sort legend entries by the latest expenditure level

```

```

sorted_legend_handles = [line for country, line in sorted(country_lines,
    ↪key=lambda x: latest_expenditure_sorted[x[0]], reverse=True)]

```

```

plt.xlabel('Year')

```

```

plt.ylabel('Expenditure_Constant_2022')

```


```

plt.title('Historical and 10-Year Random Forest Forecast for the Top 15
    ↪Countries by Expenditure_Constant_2022')

```



```
# Placing the legend outside to the left of the plot and ranked by most recent
↳expenditure level
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
↳3, 1), fontsize='small', ncol=1)
plt.grid(True, which="both", linestyle='--')
plt.show()
```



Compare model accuracy between Linear and Random

```
[22]: # Prepare a list to store the comparison of  $R^2$  scores for Linear Regression
↳and Random Forest for each country
model_comparison = []

# Calculate the  $R^2$  score for each of the top 15 countries for both Linear
↳Regression and Random Forest models
for country in top_15_countries:
    # Extract train-test data for the current country
    X_test = train_test_data[country]['X_test']
    y_test = train_test_data[country]['y_test']

    # Linear Regression  $R^2$  Score
    linear_model = models[country]
    y_pred_linear = linear_model.predict(X_test)
    r_squared_linear = r2_score(y_test, y_pred_linear)

    # Random Forest  $R^2$  Score
    rf_model = rf_models[country]
```

```

y_pred_rf = rf_model.predict(X_test)
r_squared_rf = r2_score(y_test, y_pred_rf)

# Append the comparison to the list
model_comparison.append({
    'Country': country,
    'Linear R^2 Score': r_squared_linear,
    'Random Forest R^2 Score': r_squared_rf
})

# Create a DataFrame to display the model comparison
model_comparison_df = pd.DataFrame(model_comparison)

model_comparison_df

```

```

[22]:

```

	Country	Linear R^2 Score	Random Forest R^2 Score
0	United States of America	-25.457937	-27.102278
1	China	0.694556	-5.880555
2	Russia	0.290156	-0.794920
3	Somalia	1.000000	0.000000
4	India	-0.237510	-10.619484
5	Saudi Arabia	-78.910826	-8.238869
6	United Kingdom	-0.080564	-3.584137
7	Ukraine	-0.674017	-0.720477
8	Germany	-11.194959	-7.560721
9	France	-1.866650	-2.222828
10	Japan	-2.388573	-1.605633
11	Korea, South	-3.369126	-19.531874
12	Italy	-10.028473	-2.944354
13	Australia	0.406479	-4.445582
14	Israel	-0.070533	-1.030049

```

[23]: # Prepare a list to store the summary of the Random Forest model's intercept,
      ↪ (using the mean), coefficient approximation (feature importance), and R^2
      ↪ for each country
rf_model_summary = []

# Calculate the intercept (mean value), approximate coefficient, and R^2 score
↪ for each of the top 15 countries
for country in top_15_countries:
    # Extract the Random Forest model and train-test data for the current
    ↪ country
    rf_model = rf_models[country]
    X_train = train_test_data[country]['X_train']
    y_train = train_test_data[country]['y_train']
    X_test = train_test_data[country]['X_test']
    y_test = train_test_data[country]['y_test']

```

```

# Calculate the intercept as the mean of the target value
intercept = y_train.mean()

# Approximate the "coefficient" as feature importance (since Random Forest
↳ doesn't directly provide coefficients)
feature_importance = rf_model.feature_importances_[0]

# Calculate the R2 score on the test data
y_pred = rf_model.predict(X_test)
r_squared = r2_score(y_test, y_pred)

# Append the summary to the list
rf_model_summary.append({
    'Country': country,
    'Approximate Coefficient (Feature Importance)': feature_importance,
    'Intercept (Mean Value)': intercept,
    'R2 Score': r_squared
})

# Create a DataFrame to display the model summary
rf_model_summary_df = pd.DataFrame(rf_model_summary)

rf_model_summary_df

```

```

[23]:
      Country  Approximate Coefficient (Feature Importance) \
0  United States of America                                1.0
1                China                                    1.0
2                Russia                                    1.0
3                Somalia                                    0.0
4                India                                    1.0
5            Saudi Arabia                                    1.0
6            United Kingdom                                    1.0
7                Ukraine                                    1.0
8                Germany                                    1.0
9                France                                    1.0
10               Japan                                    1.0
11           Korea, South                                    1.0
12                Italy                                    1.0
13            Australia                                    1.0
14               Israel                                    1.0

      Intercept (Mean Value)  R2 Score
0      797284.510632 -27.102278
1      134350.906012  -5.880555
2       56886.987055  -0.794920
3      126473.354300   0.000000

```

4	49049.911387	-10.619484
5	57943.995012	-8.238869
6	61806.560849	-3.584137
7	3475.918031	-0.720477
8	42074.764294	-7.560721
9	48054.356989	-2.222828
10	39764.717082	-1.605633
11	29782.014937	-19.531874
12	30678.855325	-2.944354
13	21684.055534	-4.445582
14	17716.323664	-1.030049

[48]: *# Plotting the historical data, Linear Regression forecast, and Random Forest
 ↳forecast for the top 15 countries with consistent colors*

```
plt.figure(figsize=(15, 10))

for country in sorted(top_15_countries):
    # Get the historical data for the last 10 years for the selected country
    country_data = forecast_data[forecast_data['Country'] == country]
    historical_years = country_data['Year'].tail(10)
    historical_expenditure = country_data['Expenditure_Constant_2022'].tail(10)

    # Plot the historical data with a solid line
    historical_line, = plt.plot(
        historical_years, historical_expenditure, linestyle='-',
        ↳label=f'{country} (History)'
    )

    # Plot the Linear Regression forecast for the next 10 years with a dashed
    ↳line
    linear_forecast_years = np.arange(2022, 2032)
    linear_expenditure_forecast = forecasts[country]['10_year_forecast']
    plt.plot(
        linear_forecast_years, linear_expenditure_forecast, linestyle='--',
        ↳color=historical_line.get_color(),
        label=f'{country} (Linear Forecast)'
    )

    # Plot the Random Forest forecast for the next 10 years with a dotted line
    rf_expenditure_forecast = rf_forecasts[country]['10_year_forecast']
    plt.plot(
        linear_forecast_years, rf_expenditure_forecast, linestyle=':',
        ↳color=historical_line.get_color(),
        label=f'{country} (RF Forecast)'
    )
```


```

# Sort legend entries by the latest expenditure level
sorted_legend_handles = [line for country, line in sorted(country_lines,
    ↪key=lambda x: latest_expenditure_sorted[x[0]], reverse=True)]

plt.xlabel('Year')
plt.ylabel('Expenditure_Constant_2022')
plt.title('Historical, Linear Regression Forecast, and Random Forest Forecast_
    ↪for the Top 15 Countries')

# Placing the legend outside to the left of the plot and ranked by most recent_
    ↪expenditure level
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
    ↪3, 1), fontsize='small', ncol=1)
plt.grid(True, which="both", linestyle='--')
plt.show()

```



1.1.7 Gradient Boosting Regressor Model

```

[25]: from sklearn.ensemble import GradientBoostingRegressor

# Initialize a dictionary to store the Gradient Boosting models for each country
gbr_models = {}

# Train a Gradient Boosting Regressor for each country
for country in top_15_countries:
    # Extract train-test data for the current country

```

```

X_train = train_test_data[country]['X_train']
y_train = train_test_data[country]['y_train']

# Initialize and train the Gradient Boosting Regressor
gbr_model = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1,
↳max_depth=3, random_state=42)
gbr_model.fit(X_train, y_train)

# Store the trained model
gbr_models[country] = gbr_model

# Forecast the next 10 years for each country using the trained Gradient
↳Boosting models
gbr_forecasts = {}

for country in top_15_countries:
    # Use the trained model to make predictions for the next 10 years
    gbr_model = gbr_models[country]
    forecast_10 = gbr_model.predict(forecast_years_10)

    gbr_forecasts[country] = {
        '10_year_forecast': forecast_10
    }

# Display the forecast for the first country as a sample
gbr_forecasts[sample_country]

```

```

[25]: {'10_year_forecast': array([795404.95804064, 795404.95804064, 795404.95804064,
795404.95804064,
795404.95804064, 795404.95804064, 795404.95804064, 795404.95804064,
795404.95804064, 795404.95804064])}

```

```

[49]: # Plotting the historical data, Linear Regression forecast, Random Forest
↳forecast, and Gradient Boosting Regressor forecast for the top 15 countries

plt.figure(figsize=(15, 10))

for country in sorted(top_15_countries):
    # Get the historical data for the last 10 years for the selected country
    country_data = forecast_data[forecast_data['Country'] == country]
    historical_years = country_data['Year'].tail(10)
    historical_expenditure = country_data['Expenditure_Constant_2022'].tail(10)

    # Plot the historical data with a solid line
    historical_line, = plt.plot(
        historical_years, historical_expenditure, linestyle='-',
↳label=f'{country} (History)'

```

```

)

# Plot the Linear Regression forecast for the next 10 years with a dashed
↳ line
forecast_years = np.arange(2022, 2032)
linear_expenditure_forecast = forecasts[country]['10_year_forecast']
plt.plot(
    forecast_years, linear_expenditure_forecast, linestyle='--',
↳ color=historical_line.get_color(),
    label=f'{country} (Linear Forecast)'
)

# Plot the Random Forest forecast for the next 10 years with a dotted line
rf_expenditure_forecast = rf_forecasts[country]['10_year_forecast']
plt.plot(
    forecast_years, rf_expenditure_forecast, linestyle=':',
↳ color=historical_line.get_color(),
    label=f'{country} (RF Forecast)'
)


# Plot the Gradient Boosting Regressor forecast for the next 10 years with
↳ a dash-dot line
gbr_expenditure_forecast = gbr_forecasts[country]['10_year_forecast']
plt.plot(
    forecast_years, gbr_expenditure_forecast, linestyle='-.',
↳ color=historical_line.get_color(),
    label=f'{country} (GBR Forecast)'
)

# Sort legend entries by the latest expenditure level
sorted_legend_handles = [line for country, line in sorted(country_lines,
↳ key=lambda x: latest_expenditure_sorted[x[0]], reverse=True)]

plt.xlabel('Year')
plt.ylabel('Expenditure_Constant_2022')
plt.title('Historical, Linear Regression, Random Forest, and Gradient Boosting
↳ Forecast for the Top 15 Countries')

# Placing the legend outside to the left of the plot and ranked by most recent
↳ expenditure level
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
↳ 3, 1), fontsize='small', ncol=1)
plt.grid(True, which="both", linestyle='--')
plt.show()

```



[]:

```
[50]: # Plotting the historical data, Linear Regression forecast, Random Forest
      ↪forecast, and Gradient Boosting Regressor forecast for the top 15 countries

plt.figure(figsize=(15, 10))

country_lines = [] # To store line references for sorting later

for country in sorted(top_15_countries):
    # Get the historical data for the last 10 years for the selected country
    country_data = forecast_data[forecast_data['Country'] == country]
    historical_years = country_data['Year'].tail(10)
    historical_expenditure = country_data['Expenditure_Constant_2022'].tail(10)

    # Plot the historical data with a solid line
    historical_line, = plt.plot(
        historical_years, historical_expenditure, linestyle='-', ↪
        ↪label=f'{country} (History)'
    )
    country_lines.append((country, historical_line)) # Storing line reference ↪
    ↪for legend sorting

    # Plot the Linear Regression forecast for the next 10 years with a dashed ↪
    ↪line
    forecast_years = np.arange(2022, 2032)
```



```

linear_expenditure_forecast = forecasts[country]['10_year_forecast']
plt.plot(
    forecast_years, linear_expenditure_forecast, linestyle='--',
    color=historical_line.get_color(),
    label=f'{country} (Linear Forecast)'
)

# Plot the Random Forest forecast for the next 10 years with a dotted line
rf_expenditure_forecast = rf_forecasts[country]['10_year_forecast']
plt.plot(
    forecast_years, rf_expenditure_forecast, linestyle=':',
    color=historical_line.get_color(),
    label=f'{country} (RF Forecast)'
)


# Plot the Gradient Boosting Regressor forecast for the next 10 years with
a dash-dot line
gbr_expenditure_forecast = gbr_forecasts[country]['10_year_forecast']
plt.plot(
    forecast_years, gbr_expenditure_forecast, linestyle='-.',
    color=historical_line.get_color(),
    label=f'{country} (GBR Forecast)'
)

# Sort legend entries by the latest expenditure level
sorted_legend_handles = [line for country, line in sorted(country_lines,
    key=lambda x: historical_expenditure.iloc[-1], reverse=True)]

plt.xlabel('Year')
plt.ylabel('Expenditure (Constant 2022 Dollars)')
plt.yscale('log') # Setting the y-axis to logarithmic scale
plt.title('Historical, Linear Regression, Random Forest, and Gradient Boosting
Forecast for the Top 15 Countries')

# Placing the legend outside to the left of the plot and ranked by most recent
expenditure level
plt.legend(handles=sorted_legend_handles, loc='upper left', bbox_to_anchor=(-0.
    3, 1), fontsize='small', ncol=1)
plt.grid(True, which='both', linestyle='--') # Adding gridlines for better
readability
plt.show()

```



```
[51]: # Prepare a list to store the summary of the Gradient Boosting model's  $R^2$  for
      ↪ each country
gbr_model_summary = []

# Calculate the  $R^2$  score for each of the top 15 countries for the Gradient
↪ Boosting model
for country in top_15_countries:
    # Extract train-test data for the current country
    X_test = train_test_data[country]['X_test']
    y_test = train_test_data[country]['y_test']

    # Gradient Boosting Regressor  $R^2$  Score
    gbr_model = gbr_models[country]
    y_pred_gbr = gbr_model.predict(X_test)
    r_squared_gbr = r2_score(y_test, y_pred_gbr)

    # Append the summary to the list
    gbr_model_summary.append({
        'Country': country,
        'Gradient Boosting  $R^2$  Score': r_squared_gbr
    })

# Create a DataFrame to display the Gradient Boosting model accuracy summary
gbr_model_summary_df = pd.DataFrame(gbr_model_summary)

gbr_model_summary_df
```

```
[51]:
```

	Country	Gradient Boosting R^2 Score
0	United States of America	-23.091954
1	China	-4.215515
2	Russia	-1.005995
3	Somalia	1.000000
4	India	-7.432400
5	Saudi Arabia	-10.015545
6	United Kingdom	-3.439904
7	Ukraine	-0.697479
8	Germany	-6.679121
9	France	-2.599453
10	Japan	-1.182671
11	Korea, South	-15.473089
12	Italy	-2.677620
13	Australia	-4.756625
14	Israel	-0.839813

```
[31]: # Prepare a list to store the summary of the Gradient Boosting model's intercept
      ↪ (using the mean), coefficient approximation (feature importance), and  $R^2$ 
      ↪ for each country

gbr_summary_list = [] # List to store summaries for each country

for country, gbr_model in gbr_models.items():
    # Calculate the mean of predictions as an approximation for the intercept
    mean_intercept = np.mean(gbr_model.predict(X_train)) # Assuming X_train
    ↪ contains the historical features

    # Get feature importances from the Gradient Boosting model
    feature_importance = gbr_model.feature_importances_

    # Calculate the  $R^2$  score of the model
    r_squared = gbr_model.score(X_train, y_train)

    # Prepare a dictionary for the current country's summary
    country_summary = {
        'Country': country,
        'Mean_Intercept': mean_intercept,
        'Feature_Importance': feature_importance,
        'R_squared': r_squared
    }

    # Append the dictionary to the list
    gbr_summary_list.append(country_summary)

gbr_summary_df = pd.DataFrame(gbr_summary_list)
```

gbr_summary_df

```
[31]:
```

	Country	Mean_Intercept	Feature_Importance	R_squared
0	United States of America	797284.510632	[1.0]	-94132.992031
1	China	134350.906012	[1.0]	-2626.882945
2	Russia	56886.987055	[1.0]	-277.411792
3	Somalia	126473.354300	[0.0]	-1786.684212
4	India	49049.911387	[1.0]	-163.153789
5	Saudi Arabia	57943.995012	[1.0]	-289.098608
6	United Kingdom	61806.560849	[1.0]	-296.795409
7	Ukraine	3475.918031	[1.0]	-29.979516
8	Germany	42074.764294	[1.0]	-89.737523
9	France	48054.356989	[1.0]	-138.682528
10	Japan	39764.717082	[1.0]	-73.304696
11	Korea, South	29782.014937	[1.0]	-22.886129
12	Italy	30678.855325	[1.0]	-29.425584
13	Australia	21684.055534	[1.0]	-1.814031
14	Israel	17716.323664	[1.0]	1.000000

Hyperparameter grid for Gradient Boosting Regressor

```
[34]: from sklearn.model_selection import GridSearchCV

# Define the hyperparameter grid for Gradient Boosting Regressor
param_grid = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [2, 3, 5]
}

# Select a subset of countries for demonstration purposes to save computation
↳time
sample_countries = top_15_countries[:3] # Choosing the first three countries

# Initialize a dictionary to store the best GBR models after hyperparameter
↳tuning
best_gbr_models = {}

# Hyperparameter tuning for each selected country using GridSearchCV
for country in sample_countries:
    # Extract train-test data for the current country
    X_train = train_test_data[country]['X_train']
    y_train = train_test_data[country]['y_train']

    # Initialize the GBR model
    gbr = GradientBoostingRegressor(random_state=42)

    # Set up GridSearchCV
```

```

grid_search = GridSearchCV(estimator=gbr, param_grid=param_grid, cv=3,
↪n_jobs=-1, scoring='r2', verbose=1)

# Fit the model to the training data
grid_search.fit(X_train, y_train)

# Store the best model for the current country
best_gbr_models[country] = grid_search.best_estimator_

# Display the best parameters for each of the sample countries
best_gbr_parameters = {country: best_gbr_models[country].get_params() for
↪country in sample_countries}
best_gbr_parameters

```

Fitting 3 folds for each of 27 candidates, totalling 81 fits

Fitting 3 folds for each of 27 candidates, totalling 81 fits

Fitting 3 folds for each of 27 candidates, totalling 81 fits

```

[34]: {'United States of America': {'alpha': 0.9,
    'ccp_alpha': 0.0,
    'criterion': 'friedman_mse',
    'init': None,
    'learning_rate': 0.01,
    'loss': 'squared_error',
    'max_depth': 5,
    'max_features': None,
    'max_leaf_nodes': None,
    'min_impurity_decrease': 0.0,
    'min_samples_leaf': 1,
    'min_samples_split': 2,
    'min_weight_fraction_leaf': 0.0,
    'n_estimators': 50,
    'n_iter_no_change': None,
    'random_state': 42,
    'subsample': 1.0,
    'tol': 0.0001,
    'validation_fraction': 0.1,
    'verbose': 0,
    'warm_start': False},
    'China': {'alpha': 0.9,
    'ccp_alpha': 0.0,
    'criterion': 'friedman_mse',
    'init': None,
    'learning_rate': 0.2,
    'loss': 'squared_error',
    'max_depth': 5,
    'max_features': None,

```

```

'max_leaf_nodes': None,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 200,
'n_iter_no_change': None,
'random_state': 42,
'subsample': 1.0,
'tol': 0.0001,
'validation_fraction': 0.1,
'verbose': 0,
'warm_start': False},
'Russia': {'alpha': 0.9,
'ccp_alpha': 0.0,
'criterion': 'friedman_mse',
'init': None,
'learning_rate': 0.1,
'loss': 'squared_error',
'max_depth': 2,
'max_features': None,
'max_leaf_nodes': None,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 100,
'n_iter_no_change': None,
'random_state': 42,
'subsample': 1.0,
'tol': 0.0001,
'validation_fraction': 0.1,
'verbose': 0,
'warm_start': False}}

```

```

[35]: # Feature selection and data preparation
forecast_data = data_cleaned[['Country', 'Year', 'Expenditure_Constant_2022']].
↳sort_values(by=['Country', 'Year'])

# Train-test split preparation
countries = forecast_data['Country'].unique()
train_test_data = {}
for country in countries:
    country_data = forecast_data[forecast_data['Country'] == country]
    X = country_data[['Year']]
    y = country_data['Expenditure_Constant_2022']
    train_size = int(0.8 * len(X))

```

```

X_train, X_test = X[:train_size], X[train_size:]
y_train, y_test = y[:train_size], y[train_size:]
train_test_data[country] = {
    'X_train': X_train,
    'X_test': X_test,
    'y_train': y_train,
    'y_test': y_test
}

# Redefine the sample countries and manually selected hyperparameter set
sample_countries = countries[:3]
manually_selected_params = {
    'n_estimators': 100,
    'learning_rate': 0.05,
    'max_depth': 3,
    'random_state': 42
}

# Initialize a dictionary to store the manually tuned GBR models for each
↳selected country
manually_tuned_gbr_models = {}

# Train a manually tuned Gradient Boosting Regressor for each selected country
for country in sample_countries:
    # Extract train-test data for the current country
    X_train = train_test_data[country]['X_train']
    y_train = train_test_data[country]['y_train']

    # Initialize and train the Gradient Boosting Regressor with manually
↳selected parameters
    gbr_model = GradientBoostingRegressor(**manually_selected_params)
    gbr_model.fit(X_train, y_train)

    # Store the trained model
    manually_tuned_gbr_models[country] = gbr_model

# Forecast the next 10 years for each country using the manually tuned Gradient
↳Boosting models
forecast_years_10 = pd.DataFrame({'Year': np.arange(2022, 2032)})
manually_tuned_gbr_forecasts = {}

for country in sample_countries:
    # Use the trained model to make predictions for the next 10 years
    gbr_model = manually_tuned_gbr_models[country]
    forecast_10 = gbr_model.predict(forecast_years_10)

    manually_tuned_gbr_forecasts[country] = {

```

```
    '10_year_forecast': forecast_10  
}
```

```
# Display the forecast for the first country as a sample after manual tuning  
manually_tuned_gbr_forecasts[sample_countries[0]]
```

```
[35]: {'10_year_forecast': array([176.17101111, 176.17101111, 176.17101111,  
    176.17101111,  
    176.17101111, 176.17101111, 176.17101111, 176.17101111,  
    176.17101111, 176.17101111])}
```

References

ChatGPT, (2024) GPT-4o version, OpenAI. [Large language model]. <https://chatgpt.com/>

Gelman, A., & Hill, J. (2006). Data Analysis Using Regression and Multilevel/Hierarchical Models. Cambridge University Press.

Guyon, I., & Elisseeff, A. (2003). An Introduction to Variable and Feature Selection. Journal of Machine Learning Research, 3(Mar), 1157-1182.

Hyndman, R.J., & Athanasopoulos, G. (2021) Forecasting: principles and practice, 3rd edition, OTexts: Melbourne, Australia. OTexts.com/fpp3.

LinearRegression https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html

Ridge Regression. scikit-learn. https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html

GradientBoostingRegressor <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingRegressor.html>

CountVectorizer. scikit-learn. https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html

Train-Test Split. scikit-learn. https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

[]: